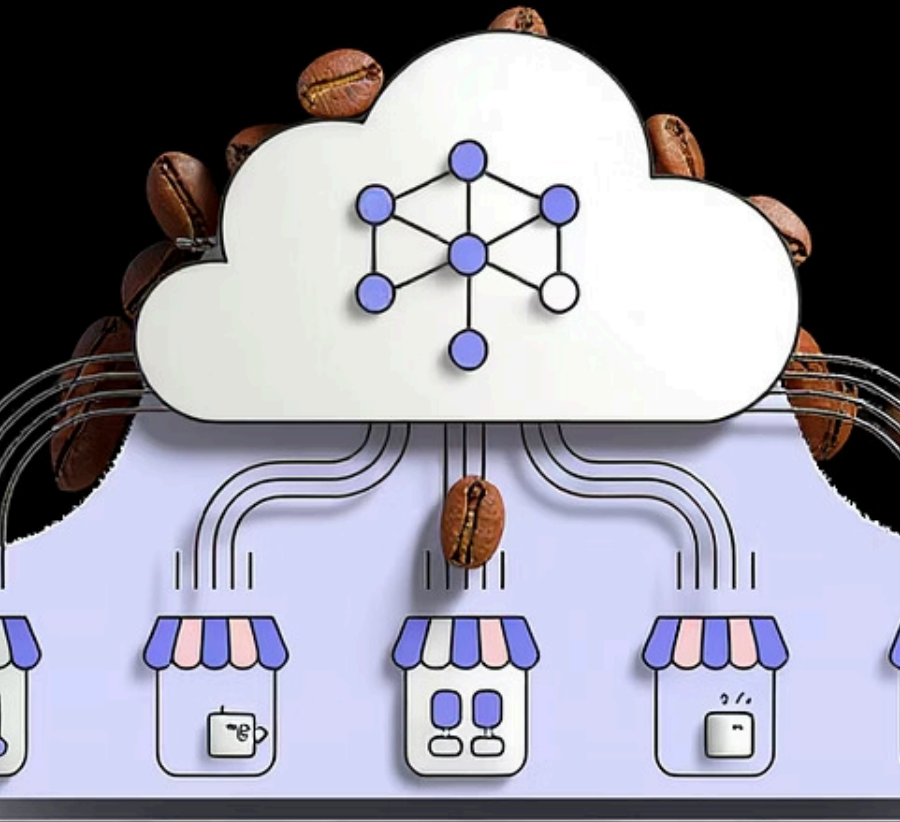




# Ana-LatteX

Unlock your cafe's potential



# Ana-LatteX – Cloud ETL & Monitoring Update

Building a scalable, cloud-based ETL pipeline for centralised analytics and operational monitoring across all café branches



Inventory optimization

Why add analytics to your data



Potential


# Project Vision

Creating a comprehensive data pipeline to transform branch operations into actionable insights:



## Branch Data Collection

CSVs from each branch uploaded nightly to **S3** buckets



## ETL Processing

**Lambda** functions process data through automated ETL pipeline



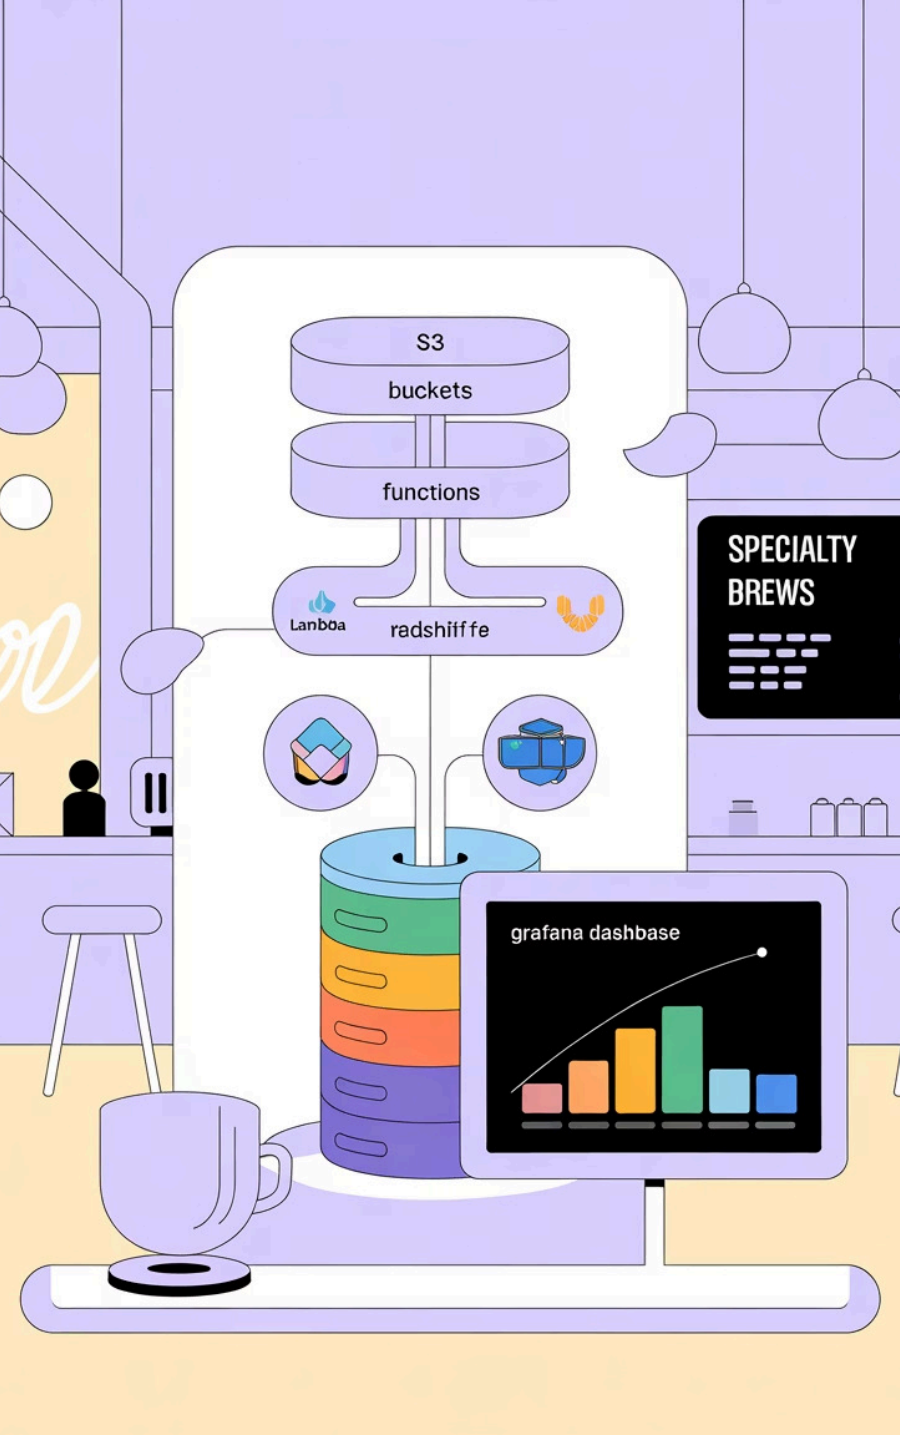
## Centralised Storage

Transformed data stored in **Redshift** for analytics and querying



## Visualisation

**Grafana** dashboards display business and operational metrics



# Sprint 2 Goals: ETL in Cloud

1

## Infrastructure as Code

Deploy cloud infrastructure using **CloudFormation** for consistent, repeatable deployments

2

## Serverless ETL

Migrate local ETL code to **AWS Lambda** for scalable, event-driven processing

3

## Automated Triggers

Configure Lambda to trigger automatically on **S3 CSV uploads** from branches

4

## Data Warehouse Integration

Connect Lambda to **Redshift** for centralised data storage and analytics

5

## Maintain PoC Pipeline

Keep **local PoC pipeline** independent for debugging and development

# Sprint 2 Deliverables: Ticket Highlights

Ticket	Deliverable	Status / Impact
1	Deployment bucket stack <i>(ana-lattex-deployment-bucket)</i>	Completed - Automates Lambda & ETL deployments
2	Lambda triggered by S3	Completed - Processes CSVs automatically; logs row count
3	ETL CloudFormation stack <i>(ana-lattex-etl-stack)</i>	Completed - Fully IaC deployment; prints "hello team"
4	Redshift Integration	Completed - Lambda connects securely in private subnet; local POC intact

# Sprint 2 Demo: Key Proof Points

## CSV Upload Trigger

Demonstrate automated Lambda activation when sample CSV is uploaded to S3 bucket

## Row Count Logging

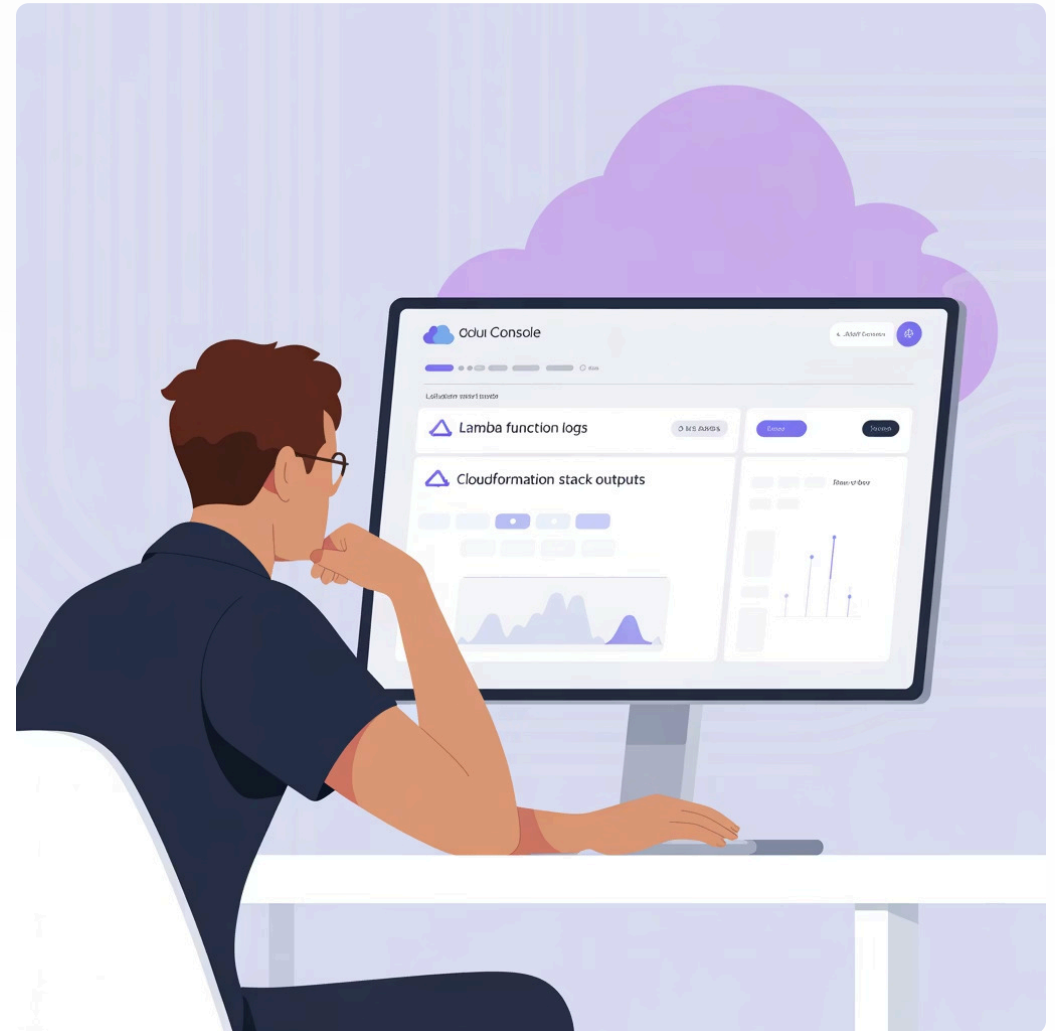
Show CloudWatch logs confirming accurate processing of uploaded data

## CloudFormation Stacks


Review ETL infrastructure deployed through IaC templates

## Local PoC Validation

Confirm local proof-of-concept pipeline remains functional



# Sprint 3 Goals: Monitoring & Visualisation



Deploy  
Grafana

Connect to  
Redshift

CloudWatch  
Integration

Business  
Dashboards

Operational  
Dashboards

- Deploy **Grafana** on EC2 using CloudFormation for consistent, repeatable setup
- Establish secure connections to both **Redshift** and **CloudWatch** data sources
- Implement version-controlled dashboards for comprehensive monitoring

# Sprint 3 Ticket Highlights



## Ticket 1: EC2 + Grafana Setup

**User Story:** Developers can visualise application and data performance

**Impact:** Secure login/password management; consistent environment



## Ticket 2: Data Source Connections

**User Story:** Data analysts can explore metrics from multiple sources

**Impact:** Dashboards backed up in GitHub; version-controlled



## Ticket 3: Business Dashboards

**User Story:** Track products sold per day, revenue per branch

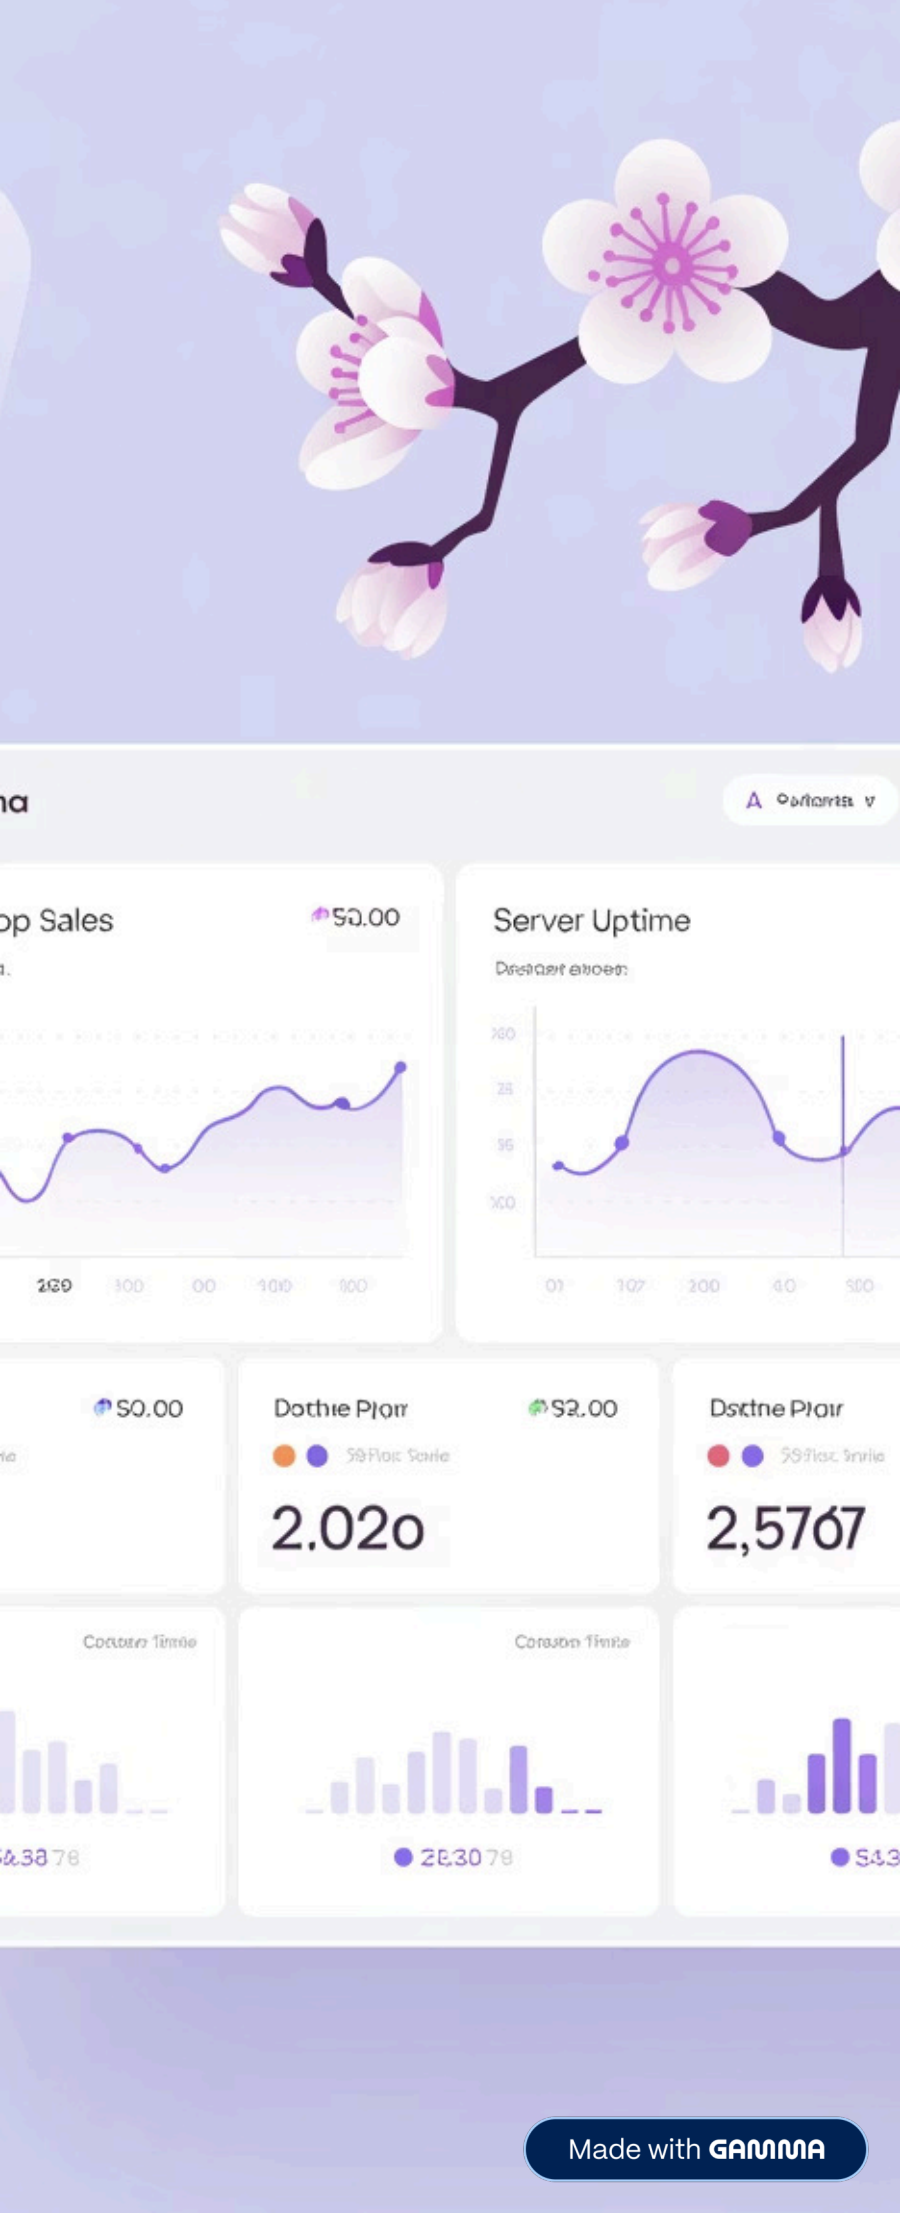
**Impact:** Data-driven decisions; branch performance comparison



## Ticket 4: Operational Dashboards

**User Story:** Monitor Lambda & EC2 performance metrics

**Impact:** Proactive maintenance; reduced downtime





# Technical Highlights



## Cloud-Ready ETL Architecture

Scalable pipeline capable of handling multiple branches with consistent processing



## Infrastructure as Code

CloudFormation deployment for Lambda, S3, EC2, Grafana ensures reproducibility



## Event-Driven Automation

S3 → Lambda → ETL → Redshift pipeline triggered automatically by data uploads



## Secure Access Controls

Redshift & Grafana access managed via SSM & Password Manager



## Development Continuity

Local PoC pipeline & unit tests remain unaffected by cloud migration



## Version-Controlled Visualisations

Grafana dashboards exportable & stored in GitHub repository

# Benefits to Business

## Centralised Insights

All branch data consolidated in Redshift for comprehensive analysis and reporting

## Operational Visibility

Real-time monitoring of Lambda & EC2 performance ensures system reliability

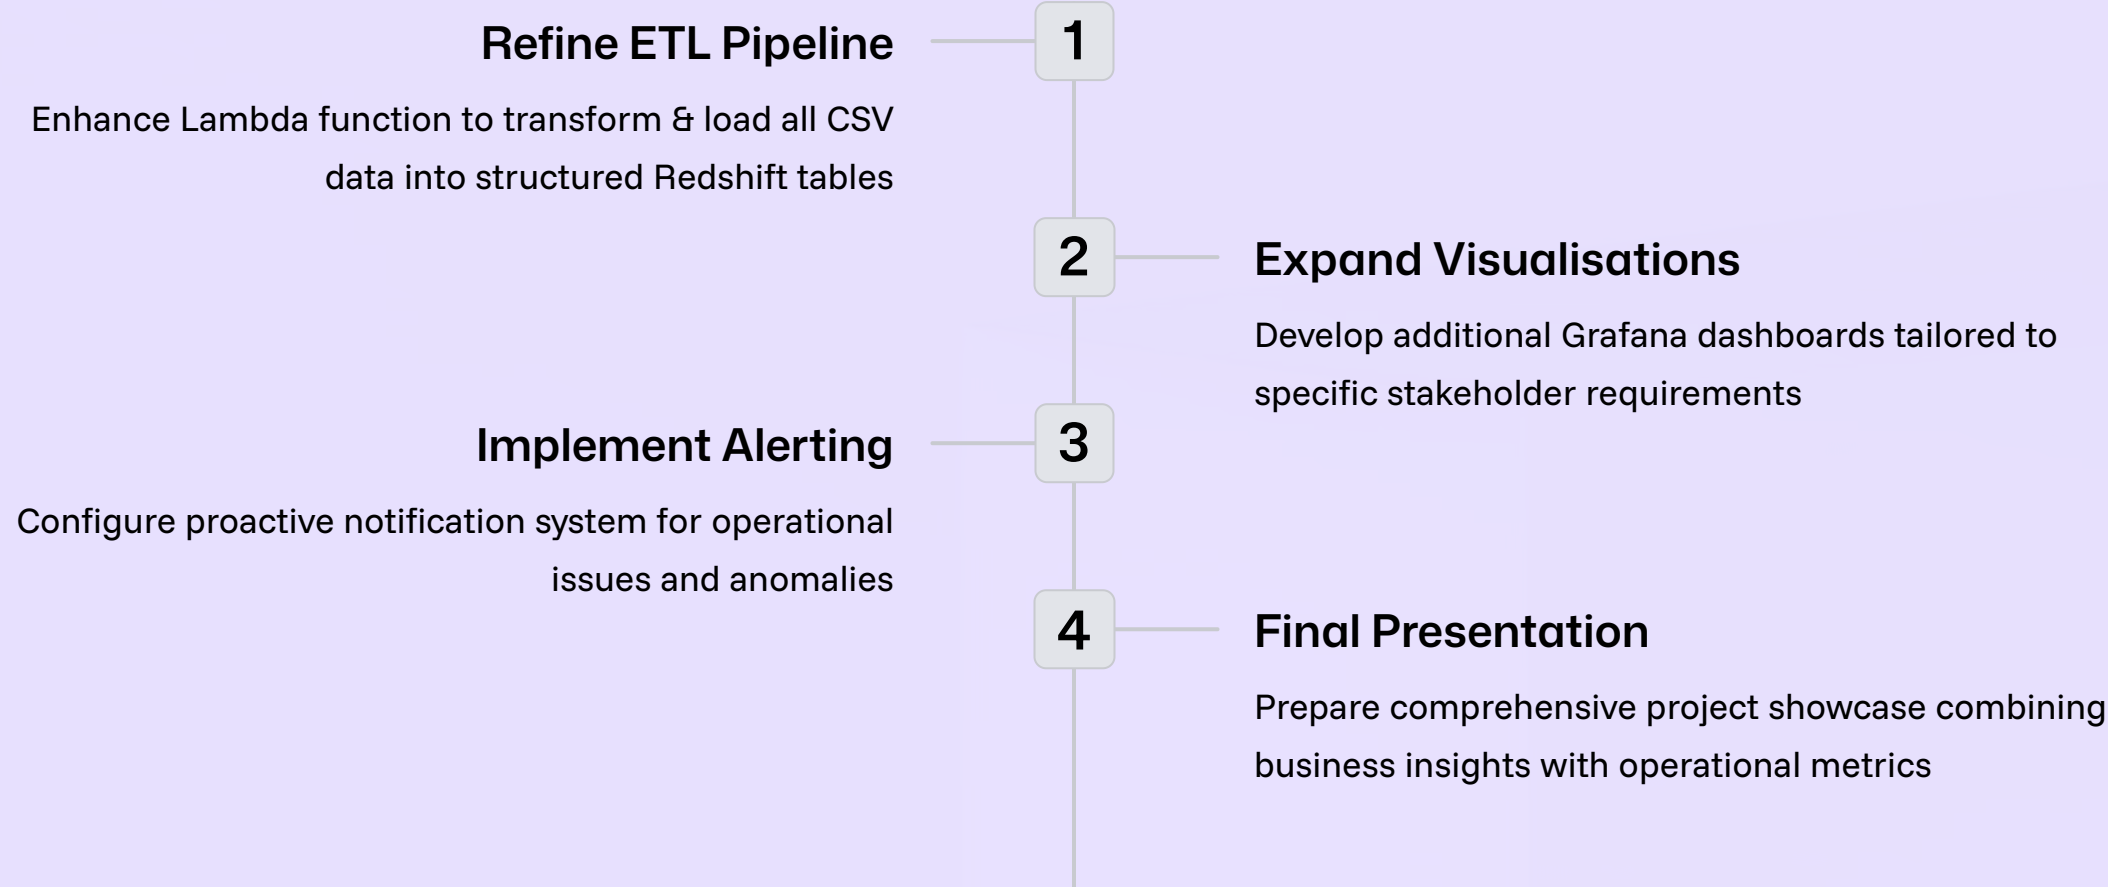
## Scalable Solution

Architecture designed to easily accommodate new branches and additional data sources

## Accelerated Decision-Making

Visual dashboards provide immediate insights into sales trends and revenue patterns

# Next Steps



📄 Our cloud ETL and monitoring solution provides the foundation for data-driven decision making across Ana-LatteX's café operations, with key components now in production.